

* HW 2 - due Thurs Sep 11

* Exam 1 on Sep 18

Plan for today : Knapsack $(1-\epsilon)$ approx, Markov Decision Process

Knapsack Problem

Given a budget $B \leftarrow$ integer

n jobs : size s_i and value v_i (assume $s_i \leq B$)
 $\{1, 2, \dots, n\}$

Goal : Find $A \subseteq [n]$ s.t.

$$\sum_{i \in A} s_i \leq B \text{ and we maximize } \sum_{i \in A} v_i$$

Thm 1 : We can solve Knapsack in time $O(n^2 \cdot \max_i v_i)$

good for small values,
but not polytime in general

Approx Algorithm

What if we don't want to assume values & want a polytime algo?

Give up on optimality.

Given a fixed error parameter ϵ (think $\epsilon = 0.01$)
 \uparrow
constant

Design an Alg with value $\geq (1-\epsilon)$ optimum and run time

$$\text{poly}(n, \log B, \sum_i \log v_i, \sum_i \log s_i)$$

$$\text{E.g. } n^{1/\epsilon} (\log B \cdot \log v_{\max})^{10}$$

Such an Alg is called Polytime Approx Scheme (PTAS), i.e., runtime
is $\text{poly}(\# \text{input bits})$ assuming ϵ is a constant

such an ϵ may be known to have $\text{poly}(\# \text{input bits})$ assuming ϵ is a constant

Thm: For knapsack, there exists a PTAS (in fact, dependency on ϵ is just $\text{poly}(1/\epsilon)$)

Plan: Replace v_i by \tilde{v}_i such we can apply pseudo-polynomial algo

Then argue the obtained soln is $(1-\epsilon)$ approx.

Proof: Define $K = \frac{\epsilon \max v_i}{n}$

$$\text{let } \tilde{v}_i := K \left\lfloor \frac{v_i}{K} \right\rfloor \Rightarrow v_i \geq \tilde{v}_i \geq v_i - K$$

Moreover, for any set S of jobs

$$\begin{aligned} \sum_{i \in S} \tilde{v}_i &\geq \sum_{i \in S} (v_i - K) \geq \sum_{i \in S} v_i - \epsilon \max v_i \\ &\geq \sum_{i \in S} v_i - \epsilon \text{opt} \quad \text{since } \text{opt} \geq \max v_i \end{aligned}$$

Alg: (1) Run Pseudo-polynomial algo with jobs having value $\frac{\tilde{v}_i}{K}$ & size s_i and budget B . (Simplified instance)

(2) Return the obtained set S of jobs

$$\begin{aligned} \text{Runtime: } O(n^2 \max v_i) &= O(n \cdot \frac{\max \tilde{v}_i}{K}) = O(n^2 \cdot \frac{\max v_i}{K}) \\ &= O(n^3/\epsilon) \end{aligned}$$

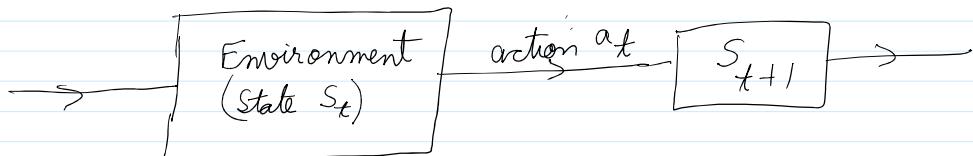
Ignores bit-complexity

Alg's value: Want to show $\sum_{i \in S} v_i \geq (1-\epsilon) \text{opt}$

Exercise on HW-2

Markov Decision Process (MDP)

Problems involving uncertainty : If Alg takes an action, the output is uncertain



Defn : • Finite state space \mathcal{S}

• Finite action space \mathcal{A}

• Start state $S_0 \in \mathcal{S}$

• Transition prob $p_a(s, s')$ for all $a \in \mathcal{A}$,
 $s, s' \in \mathcal{S}$

$$\text{where } \sum_{s' \in \mathcal{S}} p_a(s, s') = 1$$

$$\text{Input size} = O(|\mathcal{S}|^2 |\mathcal{A}|)$$

• Rewards $r_a(s) \in \mathbb{R}$

Alg (Policy) π : It tells us what action $a_t^\pi \in \mathcal{A}$ to take at t -th step given prior states.

Goal : Given a time horizon T , design an alg (policy) that plays actions to maximize $\mathbb{E}[\text{Total reward}]$

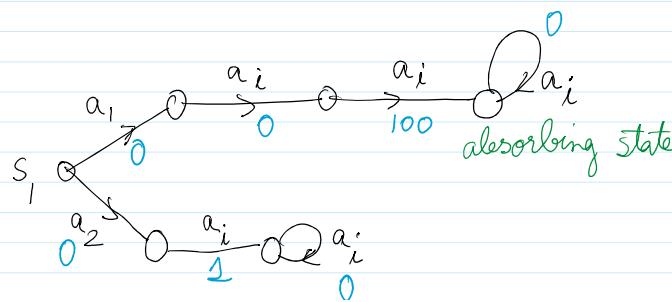
If s_t^π denotes t -th state of policy π then goal is

$$\max_{\pi} \mathbb{E} \left[\sum_{t=1}^T r_{a_t^\pi}(s_t^\pi) \right]$$

Eg 1 :

$\mathcal{S} = 6$ states

$\mathcal{A} = \{a_1, a_2\}$

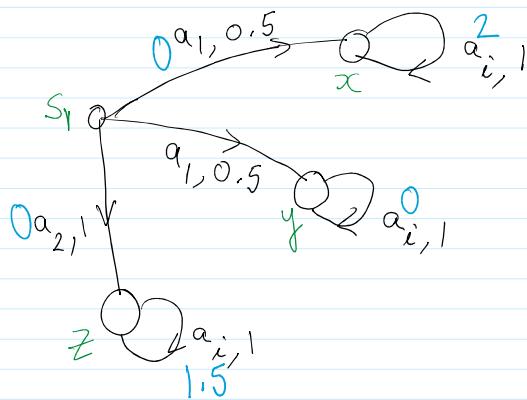


If $T=1$, what's the opt policy? a_i :: reward = 0

If $T=2$ // a_2, a_i reward = 1

If $T=3$ // a_1, a_i, a_i reward = 100

Eg 2:



$$\mathcal{S} = \{S_1, x, y, z\}$$

$$\mathcal{A} = \{a_1, a_2\}$$

If $T=1$: a_i has reward = 0

If $T=2$: $a_2, a_i \Rightarrow \text{rew} = 1.5$

$$a_1, a_i : \frac{1}{2}(0+2) + \frac{1}{2}(0+0) = 1$$

Q: what's the optimal policy given an MDP and time horizon T ?

Thm: DP can find opt policy in time $O(|\mathcal{A}| \cdot |\mathcal{S}|^2 \cdot T)$

Obs 1: We can always assume that the optimal policy is deterministic

Obs 2: Optimal action only depends on current state and remaining number of time steps.

(Does not depend on history beyond knowing current state).

Define: $\hat{V}(s, t) \leftarrow \mathbb{E} \left[\begin{array}{l} \text{Reward of optimal policy} \\ \text{Starting at state } s \in S \text{ and} \\ \text{time horizon } t \end{array} \right]$

Base case \leftarrow Easy exercise

Assume we have computed $\hat{V}(s, t')$ for $t' < t$.

$$\hat{V}^\pi(s, t) = r_{a_1^\pi}(s) + \sum_{s' \in S} p_{a_1^\pi}(s, s') * \hat{V}^\pi(s', t-1)$$

$$\hat{V}(s, t) = \max_{a \in A} \left\{ r_a(s) + \sum_{s' \in S} p_a(s, s') \cdot \hat{V}(s', t-1) \right\}$$

Solved via DP

